

Chapter 9

Cluster Analysis

Learning Objectives

After reading this chapter you should understand:

- The basic concepts of cluster analysis.
- How basic cluster algorithms work.
- How to compute simple clustering results manually.
- The different types of clustering procedures.
- The SPSS clustering outputs.

Keywords Agglomerative and divisive clustering · Chebychev distance · City-block distance · Clustering variables · Dendrogram · Distance matrix · Euclidean distance · Hierarchical and partitioning methods · Icicle diagram · k-means · Matching coefficients · Profiling clusters · Two-step clustering

Are there any market segments where Web-enabled mobile telephony is taking off in different ways? To answer this question, Okazaki (2006) applies a two-step cluster analysis by identifying segments of Internet adopters in Japan. The findings suggest that there are four clusters exhibiting distinct attitudes towards Web-enabled mobile telephony adoption. Interestingly, freelance, and highly educated professionals had the most negative perception of mobile Internet adoption, whereas clerical office workers had the most positive perception. Furthermore, housewives and company executives also exhibited a positive attitude toward mobile Internet usage. Marketing managers can now use these results to better target specific customer segments via mobile Internet services.

Introduction

Grouping similar customers and products is a fundamental marketing activity. It is used, prominently, in market segmentation. As companies cannot connect with all their customers, they have to divide markets into groups of consumers, customers, or clients (called segments) with similar needs and wants. Firms can then target each of these segments by positioning themselves in a unique segment (such as Ferrari in the high-end sports car market). While market researchers often form

market segments based on practical grounds, industry practice and wisdom, cluster analysis allows segments to be formed that are based on data that are less dependent on subjectivity.

The segmentation of customers is a standard application of cluster analysis, but it can also be used in different, sometimes rather exotic, contexts such as evaluating typical supermarket shopping paths (Larson et al. 2005) or deriving employers’ branding strategies (Moroko and Uncles 2009).

Understanding Cluster Analysis

Cluster analysis is a convenient method for identifying homogenous groups of objects called clusters. Objects (or cases, observations) in a specific cluster share many characteristics, but are very dissimilar to objects not belonging to that cluster.

Let’s try to gain a basic understanding of the cluster analysis procedure by looking at a simple example. Imagine that you are interested in segmenting your customer base in order to better target them through, for example, pricing strategies.

The first step is to decide on the characteristics that you will use to segment your customers. In other words, you have to decide which clustering variables will be included in the analysis. For example, you may want to segment a market based on customers’ price consciousness (x) and brand loyalty (y). These two variables can be measured on a 7-point scale with higher values denoting a higher degree of price consciousness and brand loyalty. The values of seven respondents are shown in Table 9.1 and the scatter plot in Fig. 9.1.

The objective of cluster analysis is to identify groups of objects (in this case, customers) that are very similar with regard to their price consciousness and brand loyalty and assign them into clusters. After having decided on the clustering variables (brand loyalty and price consciousness), we need to decide on the clustering procedure to form our groups of objects. This step is crucial for the analysis, as different procedures require different decisions prior to analysis. There is an abundance of different approaches and little guidance on which one to use in practice. We are going to discuss the most popular approaches in market research, as they can be easily computed using SPSS. These approaches are: *hierarchical methods*, *partitioning methods* (more precisely, *k-means*), and *two-step clustering*, which is largely a combination of the first two methods. Each of these procedures follows a different approach to grouping the most similar objects into a cluster and to determining each object’s cluster membership. In other words, whereas an object in a certain cluster should be as similar as possible to all the other objects in the

Table 9.1 Data

Customer	A	B	C	D	E	F	G
x	3	6	5	3	6	4	1
y	7	7	6	5	5	3	2




Fig. 9.1 Scatter plot

same cluster, it should likewise be as distinct as possible from objects in different clusters.

But how do we measure similarity? Some approaches – most notably hierarchical methods – require us to specify how similar or different objects are in order to identify different clusters. Most software packages calculate a measure of (dis)similarity by estimating the distance between pairs of objects. Objects with smaller distances between one another are more similar, whereas objects with larger distances are more dissimilar.

An important problem in the application of cluster analysis is the decision regarding how many clusters should be derived from the data. This question is explored in the next step of the analysis. Sometimes, however, we already know the number of segments that have to be derived from the data. For example, if we were asked to ascertain what characteristics distinguish frequent shoppers from infrequent ones, we need to find two different clusters. However, we do not usually know the exact number of clusters and then we face a trade-off. On the one hand, you want as few clusters as possible to make them easy to understand and actionable. On the other hand, having many clusters allows you to identify more segments and more subtle differences between segments. In an extreme case, you can address each individual separately (called one-to-one marketing) to meet consumers' varying needs in the best possible way. Examples of such a micro-marketing strategy are Puma's Mongolian Shoe BBQ (www.mongolianshoebbq.puma.com) and Nike ID (<http://nikeid.nike.com>), in which customers can fully customize a pair of shoes in a hands-on, tactile, and interactive shoe-making experience. On the other hand, the costs associated with such a strategy may be prohibitively high in many

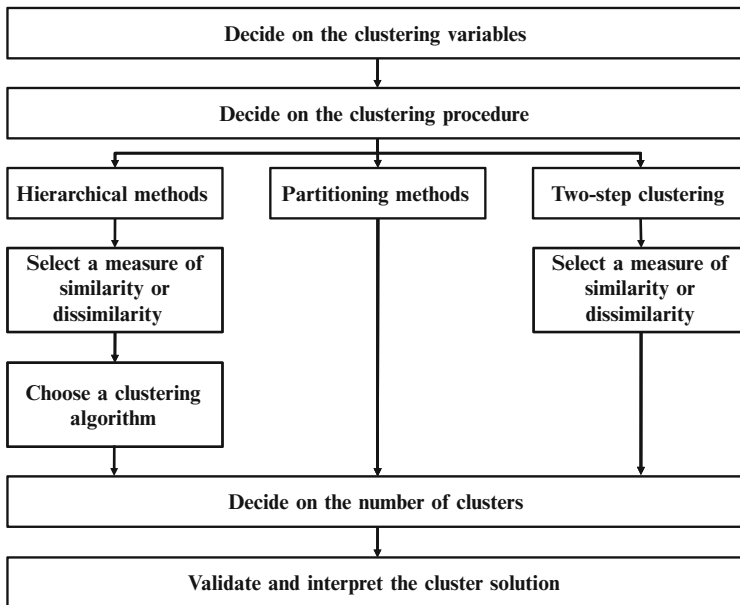


Fig. 9.2 Steps in a cluster analysis

business contexts. Thus, we have to ensure that the segments are large enough to make the targeted marketing programs profitable. Consequently, we have to cope with a certain degree of within-cluster heterogeneity, which makes targeted marketing programs less effective.

In the final step, we need to interpret the solution by defining and labeling the obtained clusters. This can be done by examining the clustering variables' mean values or by identifying explanatory variables to profile the clusters. Ultimately, managers should be able to identify customers in each segment on the basis of easily measurable variables. This final step also requires us to assess the clustering solution's stability and validity. Figure 9.2 illustrates the steps associated with a cluster analysis; we will discuss these in more detail in the following sections.

Conducting a Cluster Analysis

Decide on the Clustering Variables

At the beginning of the clustering process, we have to select appropriate variables for clustering. Even though this choice is of utmost importance, it is rarely treated as such and, instead, a mixture of intuition and data availability guide most analyses in marketing practice. However, faulty assumptions may lead to improper market

segments and, consequently, to deficient marketing strategies. Thus, great care should be taken when selecting the clustering variables.

There are several types of clustering variables and these can be classified into *general* (independent of products, services or circumstances) and *specific* (related to both the customer and the product, service and/or particular circumstance), on the one hand, and *observable* (i.e., measured directly) and *unobservable* (i.e., inferred) on the other. Table 9.2 provides several types and examples of clustering variables.

Table 9.2 Types and examples of clustering variables

	General	Specific
Observable (directly measurable)	Cultural, geographic, demographic, socio-economic	User status, usage frequency, store and brand loyalty
Unobservable (inferred)	Psychographics, values, personality, lifestyle	Benefits, perceptions, attitudes, intentions, preferences

Adapted from Wedel and Kamakura (2000)

The types of variables used for cluster analysis provide different segments and, thereby, influence segment-targeting strategies. Over the last decades, attention has shifted from more traditional general clustering variables towards product-specific unobservable variables. The latter generally provide better guidance for decisions on marketing instruments’ effective specification. It is generally acknowledged that segments identified by means of specific unobservable variables are usually more homogenous and their consumers respond consistently to marketing actions (see Wedel and Kamakura 2000). However, consumers in these segments are also frequently hard to identify from variables that are easily measured, such as demographics. Conversely, segments determined by means of generally observable variables usually stand out due to their identifiability but often lack a unique response structure.¹ Consequently, researchers often combine different variables (e.g., multiple lifestyle characteristics combined with demographic variables), benefiting from each ones strengths.

In some cases, the choice of clustering variables is apparent from the nature of the task at hand. For example, a managerial problem regarding corporate communications will have a fairly well defined set of clustering variables, including contenders such as awareness, attitudes, perceptions, and media habits. However, this is not always the case and researchers have to choose from a set of candidate variables.

Whichever clustering variables are chosen, it is important to select those that provide a clear-cut differentiation between the segments regarding a specific managerial objective.² More precisely, criterion validity is of special interest; that is, the extent to which the “independent” clustering variables are associated with

¹See Wedel and Kamakura (2000).

²Tonks (2009) provides a discussion of segment design and the choice of clustering variables in consumer markets.

one or more “dependent” variables not included in the analysis. Given this relationship, there should be significant differences between the “dependent” variable(s) across the clusters. These associations may or may not be causal, but it is essential that the clustering variables distinguish the “dependent” variable(s) significantly. Criterion variables usually relate to some aspect of behavior, such as purchase intention or usage frequency.

Generally, you should avoid using an abundance of clustering variables, as this increases the odds that the variables are no longer dissimilar. If there is a high degree of collinearity between the variables, they are not sufficiently unique to identify distinct market segments. If highly correlated variables are used for cluster analysis, specific aspects covered by these variables will be overrepresented in the clustering solution. In this regard, absolute correlations above 0.90 are always problematic. For example, if we were to add another variable called *brand preference* to our analysis, it would virtually cover the same aspect as *brand loyalty*. Thus, the concept of being attached to a brand would be overrepresented in the analysis because the clustering procedure does not differentiate between the clustering variables in a conceptual sense. Researchers frequently handle this issue by applying cluster analysis to the observations’ factor scores derived from a previously carried out factor analysis. However, according to Dolnicar and Grün (2009), this factor-cluster segmentation approach can lead to several problems:

1. The data are pre-processed and the clusters are identified on the basis of transformed values, not on the original information, which leads to different results.
2. In factor analysis, the factor solution does not explain a certain amount of variance; thus, information is discarded before segments have been identified or constructed.
3. Eliminating variables with low loadings on all the extracted factors means that, potentially, the most important pieces of information for the identification of niche segments are discarded, making it impossible to ever identify such groups.
4. The interpretations of clusters based on the original variables become questionable given that the segments have been constructed using factor scores.

Several studies have shown that the factor-cluster segmentation significantly reduces the success of segment recovery.³ Consequently, you should rather reduce the number of items in the questionnaire’s pre-testing phase, retaining a reasonable number of relevant, non-redundant questions that you believe differentiate the segments well. However, if you have your doubts about the data structure, factor-clustering segmentation may still be a better option than discarding items that may conceptually be necessary.

Furthermore, we should keep the sample size in mind. First and foremost, this relates to issues of managerial relevance as segments’ sizes need to be substantial to ensure that targeted marketing programs are profitable. From a statistical perspective, every additional variable requires an over-proportional increase in

³See the studies by Arabie and Hubert (1994), Sheppard (1996), or Dolnicar and Grün (2009).

observations to ensure valid results. Unfortunately, there is no generally accepted rule of thumb regarding minimum sample sizes or the relationship between the objects and the number of clustering variables used.

In a related methodological context, Formann (1984) recommends a sample size of at least 2^m , where m equals the number of clustering variables. This can only provide rough guidance; nevertheless, we should pay attention to the relationship between the objects and clustering variables. It does not, for example, appear logical to cluster ten objects using ten variables. Keep in mind that no matter how many variables are used and no matter how small the sample size, cluster analysis will always render a result!

Ultimately, the choice of clustering variables always depends on contextual influences such as data availability or resources to acquire additional data. Marketing researchers often overlook the fact that the choice of clustering variables is closely connected to data quality. Only those variables that ensure that high quality data can be used should be included in the analysis. This is very important if a segmentation solution has to be managerially useful. Furthermore, data are of high quality if the questions asked have a strong theoretical basis, are not contaminated by respondent fatigue or response styles, are recent, and thus reflect the current market situation (Dolnicar and Lazarevski 2009). Lastly, the requirements of other managerial functions within the organization often play a major role. Sales and distribution may as well have a major influence on the design of market segments. Consequently, we have to be aware that subjectivity and common sense agreement will (and should) always impact the choice of clustering variables.

Decide on the Clustering Procedure

By choosing a specific clustering procedure, we determine how clusters are to be formed. This always involves optimizing some kind of criterion, such as minimizing the within-cluster variance (i.e., the clustering variables' overall variance of objects in a specific cluster), or maximizing the distance between the objects or clusters. The procedure could also address the question of how to determine the (dis)similarity between objects in a newly formed cluster and the remaining objects in the dataset.

There are many different clustering procedures and also many ways of classifying these (e.g., overlapping versus non-overlapping, unimodal versus multimodal, exhaustive versus non-exhaustive).⁴ A practical distinction is the differentiation between *hierarchical* and *partitioning methods* (most notably the *k-means* procedure), which we are going to discuss in the next sections. We also introduce *two-step clustering*, which combines the principles of hierarchical and partitioning methods and which has recently gained increasing attention from market research practice.

⁴See Wedel and Kamakura (2000), Dolnicar (2003), and Kaufman and Rousseeuw (2005) for a review of clustering techniques.

Hierarchical Methods

Hierarchical clustering procedures are characterized by the tree-like structure established in the course of the analysis. Most hierarchical techniques fall into a category called *agglomerative clustering*. In this category, clusters are consecutively formed from objects. Initially, this type of procedure starts with each object representing an individual cluster. These clusters are then sequentially merged according to their similarity. First, the two most similar clusters (i.e., those with the smallest distance between them) are merged to form a new cluster at the bottom of the hierarchy. In the next step, another pair of clusters is merged and linked to a higher level of the hierarchy, and so on. This allows a hierarchy of clusters to be established from the bottom up. In Fig. 9.3 (left-hand side), we show how agglomerative clustering assigns additional objects to clusters as the cluster size increases.

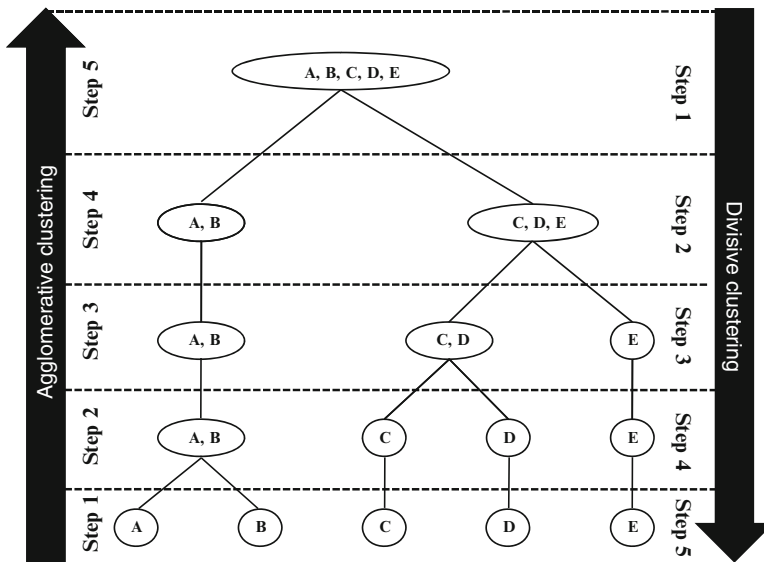


Fig. 9.3 Agglomerative and divisive clustering

A cluster hierarchy can also be generated top-down. In this *divisive clustering*, all objects are initially merged into a single cluster, which is then gradually split up. Figure 9.3 illustrates this concept (right-hand side). As we can see, in both agglomerative and divisive clustering, a cluster on a higher level of the hierarchy always encompasses all clusters from a lower level. This means that if an object is assigned to a certain cluster, there is no possibility of reassigning this object to another cluster. This is an important distinction between these types of clustering and partitioning methods such as k-means, which we will explore in the next section.

Divisive procedures are quite rarely used in market research. We therefore concentrate on the agglomerative clustering procedures. There are various types

of agglomerative procedures. However, before we discuss these, we need to define how similarities or dissimilarities are measured between pairs of objects.

Select a Measure of Similarity or Dissimilarity

There are various measures to express (dis)similarity between pairs of objects. A straightforward way to assess two objects' proximity is by drawing a straight line between them. For example, when we look at the scatter plot in Fig. 9.1, we can easily see that the length of the line connecting observations B and C is much shorter than the line connecting B and G. This type of distance is also referred to as *Euclidean distance* (or straight-line distance) and is the most commonly used type when it comes to analyzing ratio or interval-scaled data.⁵ In our example, we have ordinal data, but market researchers usually treat ordinal data as metric data to calculate distance metrics by assuming that the scale steps are equidistant (very much like in factor analysis, which we discussed in Chap. 8). To use a hierarchical clustering procedure, we need to express these distances mathematically. By taking the data in Table 9.1 into consideration, we can easily compute the Euclidean distance between customer B and customer C (generally referred to as $d(B,C)$) with regard to the two variables x and y by using the following formula:

$$d_{Euclidean}(B, C) = \sqrt{(x_B - x_C)^2 + (y_B - y_C)^2}$$

The Euclidean distance is the square root of the sum of the squared differences in the variables' values. Using the data from Table 9.1, we obtain the following:

$$d_{Euclidean}(B, C) = \sqrt{(6 - 5)^2 + (7 - 6)^2} = \sqrt{2} = 1.414$$

This distance corresponds to the length of the line that connects objects B and C. In this case, we only used two variables but we can easily add more under the root sign in the formula. However, each additional variable will add a dimension to our research problem (e.g., with six clustering variables, we have to deal with six dimensions), making it impossible to represent the solution graphically. Similarly, we can compute the distance between customer B and G, which yields the following:

$$d_{Euclidean}(B, G) = \sqrt{(6 - 1)^2 + (7 - 2)^2} = \sqrt{50} = 7.071$$

Likewise, we can compute the distance between all other pairs of objects. All these distances are usually expressed by means of a *distance matrix*. In this distance matrix, the non-diagonal elements express the distances between pairs of objects

⁵Note that researchers also often use the squared Euclidean distance.

and zeros on the diagonal (the distance from each object to itself is, of course, 0). In our example, the distance matrix is an 8×8 table with the lines and rows representing the objects (i.e., customers) under consideration (see Table 9.3). As the distance between objects B and C (in this case 1.414 units) is the same as between C and B, the distance matrix is symmetrical. Furthermore, since the distance between an object and itself is zero, one need only look at either the lower or upper non-diagonal elements.

Table 9.3 Euclidean distance matrix

Objects	A	B	C	D	E	F	G
A	0						
B	3	0					
C	2.236	1.414	0				
D	2	3.606	2.236	0			
E	3.606	2	1.414	3	0		
F	4.123	4.472	3.162	2.236	2.828	0	
G	5.385	7.071	5.657	3.606	5.831	3.162	0

There are also alternative distance measures: The *city-block distance* uses the sum of the variables' absolute differences. This is often called the Manhattan metric as it is akin to the walking distance between two points in a city like New York's Manhattan district, where the distance equals the number of blocks in the directions North-South and East-West. Using the city-block distance to compute the distance between customers B and C (or C and B) yields the following:

$$d_{\text{City-block}}(B, C) = |x_B - x_C| + |y_B - y_C| = |6 - 5| + |7 - 6| = 2$$

The resulting distance matrix is in Table 9.4.

Table 9.4 City-block distance matrix

Objects	A	B	C	D	E	F	G
A	0						
B	3	0					
C	3	2	0				
D	2	5	3	0			
E	5	2	2	3	0		
F	5	6	4	3	4	0	
G	7	10	8	5	8	4	0

Lastly, when working with metric (or ordinal) data, researchers frequently use the *Chebychev distance*, which is the maximum of the absolute difference in the clustering variables' values. In respect of customers B and C, this result is:

$$d_{\text{Chebychev}}(B, C) = \max(|x_B - x_C|, |y_B - y_C|) = \max(|6 - 5|, |7 - 6|) = 1$$

Figure 9.4 illustrates the interrelation between these three distance measures regarding two objects, C and G, from our example.




Fig. 9.4 Distance measures

There are other distance measures such as the Angular, Canberra or Mahalanobis distance. In many situations, the latter is desirable as it compensates for collinearity between the clustering variables. However, it is (unfortunately) not menu-accessible in SPSS.

In many analysis tasks, the variables under consideration are measured on different scales or levels. This would be the case if we extended our set of clustering variables by adding another ordinal variable representing the customers' income measured by means of, for example, 15 categories. Since the absolute variation of the income variable would be much greater than the variation of the remaining two variables (remember, that x and y are measured on 7-point scales), this would clearly distort our analysis results. We can resolve this problem by standardizing the data prior to the analysis.

Different standardization methods are available, such as the simple z standardization, which rescales each variable to have a mean of 0 and a standard deviation of 1 (see Chap. 5). In most situations, however, standardization by range (e.g., to a range of 0 to 1 or -1 to 1) performs better.⁶ We recommend standardizing the data in general, even though this procedure can reduce or inflate the variables' influence on the clustering solution.

⁶See Milligan and Cooper (1988).

Another way of (implicitly) standardizing the data is by using the correlation between the objects instead of distance measures. For example, suppose a respondent rated price consciousness 2 and brand loyalty 3. Now suppose a second respondent indicated 5 and 6, whereas a third rated these variables 3 and 3. Euclidean, city-block, and Chebychev distances would indicate that the first respondent is more similar to the third than to the second. Nevertheless, one could convincingly argue that the first respondent’s ratings are more similar to the second’s, as both rate brand loyalty higher than price consciousness. This can be accounted for by computing the correlation between two vectors of values as a measure of similarity (i.e., high correlation coefficients indicate a high degree of similarity). Consequently, similarity is no longer defined by means of the difference between the answer categories but by means of the similarity of the answering profiles. Using correlation is also a way of standardizing the data implicitly.

Whether you use correlation or one of the distance measures depends on whether you think the relative magnitude of the variables within an object (which favors correlation) matters more than the relative magnitude of each variable across objects (which favors distance). However, it is generally recommended that one uses correlations when applying clustering procedures that are susceptible to outliers, such as complete linkage, average linkage or centroid (see next section).

Whereas the distance measures presented thus far can be used for metrically and – in general – ordinally scaled data, applying them to nominal or binary data is meaningless. In this type of analysis, you should rather select a similarity measure expressing the degree to which variables’ values share the same category. These so-called *matching coefficients* can take different forms but rely on the same allocation scheme shown in Table 9.5.

Table 9.5 Allocation scheme for matching coefficients

		Object 1	
		Number of variables with category 1	Number of variables with category 2
Object 2	Number of variables with category 1	a	b
	Number of variables with category 2	c	d

Based on the allocation scheme in Table 9.5, we can compute different matching coefficients, such as the simple matching coefficient (SM):

$$SM = \frac{a + d}{a + b + c + d}$$

This coefficient is useful when both positive and negative values carry an equal degree of information. For example, gender is a symmetrical attribute because the number of males and females provides an equal degree of information.

Let's take a look at an example by assuming that we have a dataset with three binary variables: gender (male = 1, female = 2), customer (customer = 1, non-customer = 2), and disposable income (low = 1, high = 2). The first object is a male non-customer with a high disposable income, whereas the second object is a female non-customer with a high disposable income. According to the scheme in Table 9.4, $a = b = 0$, $c = 1$ and $d = 2$, with the simple matching coefficient taking a value of 0.667.

Two other types of matching coefficients, which do not equate the joint absence of a characteristic with similarity and may, therefore, be of more value in segmentation studies, are the *Jaccard (JC)* and the *Russel and Rao (RR)* coefficients. They are defined as follows:

$$JC = \frac{a}{a + b + c}$$

$$RR = \frac{a}{a + b + c + d}$$

These matching coefficients are – just like the distance measures – used to determine a cluster solution. There are many other matching coefficients such as Yule's Q, Kulczynski or Ochiai, but since most applications of cluster analysis rely on metric or ordinal data, we will not discuss these in greater detail.⁷

For nominal variables with more than two categories, you should always convert the categorical variable into a set of binary variables in order to use matching coefficients. When you have ordinal data, you should always use distance measures such as Euclidean distance. Even though using matching coefficients would be feasible and – from a strictly statistical standpoint – even more appropriate, you would disregard variable information in the sequence of the categories. In the end, a respondent who indicates that he or she is very loyal to a brand is going to be closer to someone who is somewhat loyal than a respondent who is not loyal at all. Furthermore, distance measures best represent the concept of proximity, which is fundamental to cluster analysis.

Most datasets contain variables that are measured on multiple scales. For example, a market research questionnaire may ask about the respondent's income, product ratings, and last brand purchased. Thus, we have to consider variables measured on a ratio, ordinal, and nominal scale. How can we simultaneously incorporate these variables into one analysis? Unfortunately, this problem cannot be easily resolved and, in fact, many market researchers simply ignore the scale level. Instead, they use one of the distance measures discussed in the context of metric (and ordinal) data. Even though this approach may slightly change the results when compared to those using matching coefficients, it should not be rejected. Cluster analysis is mostly an exploratory technique whose results provide a rough guidance for managerial decisions. Despite this, there are several procedures that allow a simultaneous integration of these variables into one analysis.

⁷See Wedel and Kamakura (2000) for more information on alternative matching coefficients.

First, we could compute distinct distance matrices for each group of variables; that is, one distance matrix based on, for example, ordinally scaled variables and another based on nominal variables. Afterwards, we can simply compute the weighted arithmetic mean of the distances and use this average distance matrix as the input for the cluster analysis. However, the weights have to be determined a priori and improper weights may result in a biased treatment of different variable types. Furthermore, the computation and handling of distance matrices are not trivial. Using the SPSS syntax, one has to manually add the `MATRIX` subcommand, which exports the initial distance matrix into a new data file. Go to the [Web Appendix](#) (→ Chap. 5) to learn how to modify the SPSS syntax accordingly.

Second, we could dichotomize all variables and apply the matching coefficients discussed above. In the case of metric variables, this would involve specifying categories (e.g., low, medium, and high income) and converting these into sets of binary variables. In most cases, however, the specification of categories would be rather arbitrary and, as mentioned earlier, this procedure could lead to a severe loss of information.

In the light of these issues, you should avoid combining metric and nominal variables in a single cluster analysis, but if this is not feasible, the *two-step clustering procedure* provides a valuable alternative, which we will discuss later. Lastly, the choice of the (dis)similarity measure is not extremely critical to recovering the underlying cluster structure. In this regard, the choice of the clustering algorithm is far more important. We therefore deal with this aspect in the following section.

Select a Clustering Algorithm

After having chosen the distance or similarity measure, we need to decide which clustering algorithm to apply. There are several agglomerative procedures and they can be distinguished by the way they define the distance from a newly formed cluster to a certain object, or to other clusters in the solution. The most popular agglomerative clustering procedures include the following:

- *Single linkage* (nearest neighbor): The distance between two clusters corresponds to the shortest distance between any two members in the two clusters.
- *Complete linkage* (furthest neighbor): The oppositional approach to single linkage assumes that the distance between two clusters is based on the longest distance between any two members in the two clusters.
- *Average linkage*: The distance between two clusters is defined as the average distance between all pairs of the two clusters' members.
- *Centroid*: In this approach, the geometric center (centroid) of each cluster is computed first. The distance between the two clusters equals the distance between the two centroids.

Figures 9.5–9.8 illustrate these linkage procedures for two randomly framed clusters.

Fig. 9.5 Single linkage




Fig. 9.6 Complete linkage




Fig. 9.7 Average linkage





Fig. 9.8 Centroid



Each of these linkage algorithms can yield totally different results when used on the same dataset, as each has its specific properties. As the single linkage algorithm is based on minimum distances, it tends to form one large cluster with the other clusters containing only one or few objects each. We can make use of this “chaining effect” to detect outliers, as these will be merged with the remaining objects – usually at very large distances – in the last steps of the analysis. Generally, single linkage is considered the most versatile algorithm. Conversely, the complete linkage method is strongly affected by outliers, as it is based on maximum distances. Clusters produced by this method are likely to be rather compact and tightly clustered. The average linkage and centroid algorithms tend to produce clusters with rather low within-cluster variance and similar sizes. However, both procedures are affected by outliers, though not as much as complete linkage.

Another commonly used approach in hierarchical clustering is *Ward's method*. This approach does not combine the two most similar objects successively. Instead, those objects whose merger increases the overall within-cluster variance to the smallest possible degree, are combined. If you expect somewhat equally sized clusters and the dataset does not include outliers, you should always use Ward's method.

To better understand how a clustering algorithm works, let's manually examine some of the single linkage procedure's calculation steps. We start off by looking at the initial (Euclidean) distance matrix in Table 9.3. In the very first step, the two objects exhibiting the smallest distance in the matrix are merged. Note that we always merge those objects with the smallest distance, regardless of the clustering procedure (e.g., single or complete linkage). As we can see, this happens to two pairs of objects, namely B and C ($d(B, C) = 1.414$), as well as C and E ($d(C, E) = 1.414$). In the next step, we will see that it does not make any difference whether we first merge the one or the other, so let's proceed by forming a new cluster, using objects B and C.

Having made this decision, we then form a new distance matrix by considering the single linkage decision rule as discussed above. According to this rule, the distance from, for example, object A to the newly formed cluster is the minimum of $d(A, B)$ and $d(A, C)$. As $d(A, C)$ is smaller than $d(A, B)$, the distance from A to the newly formed cluster is equal to $d(A, C)$; that is, 2.236. We also compute the distances from cluster [B,C] (clusters are indicated by means of squared brackets) to all other objects (i.e. D, E, F, G) and simply copy the remaining distances – such as $d(E, F)$ – that the previous clustering has not affected. This yields the distance matrix shown in Table 9.6.

Continuing the clustering procedure, we simply repeat the last step by merging the objects in the new distance matrix that exhibit the smallest distance (in this case, the newly formed cluster [B, C] and object E) and calculate the distance from this cluster to all other objects. The result of this step is described in Table 9.7.

Try to calculate the remaining steps yourself and compare your solution with the distance matrices in the following Tables 9.8–9.10.

Table 9.6 Distance matrix after first clustering step (single linkage)

Objects	A	B, C	D	E	F	G
A	0					
B, C	2.236	0				
D	2	2.236	0			
E	3.606	1.414	3	0		
F	4.123	3.162	2.236	2.828	0	
G	5.385	5.657	3.606	5.831	3.162	0

Table 9.7 Distance matrix after second clustering step (single linkage)

Objects	A	B, C, E	D	F	G
A	0				
B, C, E	2.236	0			
D	2	2.236	0		
F	4.123	2.828	2.236	0	
G	5.385	5.657	3.606	3.162	0

Table 9.8 Distance matrix after third clustering step (single linkage)

Objects	A, D	B, C, E	F	G
A, D	0			
B, C, E	2.236	0		
F	2.236	2.828	0	
G	3.606	5.657	3.162	0

Table 9.9 Distance matrix after fourth clustering step (single linkage)

Objects	A, B, C, D, E	F	G
A, B, C, D, E	0		
F	2.236	0	
G	3.606	3.162	0

Table 9.10 Distance matrix after fifth clustering step (single linkage)

Objects	A, B, C, D, E, F	G
A, B, C, D, E, F	0	
G	3.162	0

By following the single linkage procedure, the last steps involve the merger of cluster [A,B,C,D,E,F] and object G at a distance of 3.162. Do you get the same results? As you can see, conducting a basic cluster analysis manually is not that hard at all – not if there are only a few objects in the dataset.

A common way to visualize the cluster analysis's progress is by drawing a *dendrogram*, which displays the distance level at which there was a combination of objects and clusters (Fig. 9.9).

We read the dendrogram from left to right to see at which distance objects have been combined. For example, according to our calculations above, objects B, C, and E are combined at a distance level of 1.414.




Fig. 9.9 Dendrogram

Decide on the Number of Clusters

An important question we haven't yet addressed is how to decide on the number of clusters to retain from the data. Unfortunately, hierarchical methods provide only very limited guidance for making this decision. The only meaningful indicator relates to the distances at which the objects are combined. Similar to factor analysis's scree plot, we can seek a solution in which an additional combination of clusters or objects would occur at a greatly increased distance. This raises the issue of what a great distance is, of course.

One potential way to solve this problem is to plot the number of clusters on the x-axis (starting with the one-cluster solution at the very left) against the distance at which objects or clusters are combined on the y-axis. Using this plot, we then search for the distinctive break (*elbow*). SPSS does not produce this plot automatically – you have to use the distances provided by SPSS to draw a line chart by using a common spreadsheet program such as Microsoft Excel.

Alternatively, we can make use of the dendrogram which essentially carries the same information. SPSS provides a dendrogram; however, this differs slightly from the one presented in Fig. 9.9. Specifically, SPSS rescales the distances to a range of 0–25; that is, the last merging step to a one-cluster solution takes place at a (rescaled) distance of 25. The rescaling often lengthens the merging steps, thus making breaks occurring at a greatly increased distance level more obvious.

Despite this, this distance-based decision rule does not work very well in all cases. It is often difficult to identify where the break actually occurs. This is also the case in our example above. By looking at the dendrogram, we could justify a two-cluster solution ([A,B,C,D,E,F] and [G]), as well as a five-cluster solution ([B,C,E], [A], [D], [F], [G]).

Research has suggested several other procedures for determining the number of clusters in a dataset. Most notably, the *variance ratio criterion* (VRC) by Calinski and Harabasz (1974) has proven to work well in many situations.⁸ For a solution with n objects and k segments, the criterion is given by:

$$VRC_k = (SS_B / (k - 1)) / (SS_W / (n - k)),$$

where SS_B is the sum of the squares between the segments and SS_W is the sum of the squares within the segments. The criterion should seem familiar, as this is nothing but the F-value of a one-way ANOVA, with k representing the factor levels. Consequently, the VRC can easily be computed using SPSS, even though it is not readily available in the clustering procedures' outputs.

To finally determine the appropriate number of segments, we compute ω_k for each segment solution as follows:

$$\omega_k = (VRC_{k+1} - VRC_k) - (VRC_k - VRC_{k-1}).$$

In the next step, we choose the number of segments k that minimizes the value in ω_k . Owing to the term VRC_{k-1} , the minimum number of clusters that can be selected is three, which is a clear disadvantage of the criterion, thus limiting its application in practice.

Overall, the data can often only provide rough guidance regarding the number of clusters you should select; consequently, you should rather revert to practical considerations. Occasionally, you might have a priori knowledge, or a theory on which you can base your choice. However, first and foremost, you should ensure that your results are interpretable and meaningful. Not only must the number of clusters be small enough to ensure manageability, but each segment should also be large enough to warrant strategic attention.

Partitioning Methods: k-means

Another important group of clustering procedures are partitioning methods. As with hierarchical clustering, there is a wide array of different algorithms; of these, the *k-means procedure* is the most important one for market research.⁹ The k-means algorithm follows an entirely different concept than the hierarchical methods discussed before. This algorithm is not based on distance measures such as Euclidean distance or city-block distance, but uses the within-cluster variation as a

⁸Milligan and Cooper (1985) compare various criteria.

⁹Note that the k-means algorithm is one of the simplest non-hierarchical clustering methods. Several extensions, such as k-medoids (Kaufman and Rousseeuw 2005) have been proposed to handle problematic aspects of the procedure. More advanced methods include finite mixture models (McLachlan and Peel 2000), neural networks (Bishop 2006), and self-organizing maps (Kohonen 1982). Andrews and Currim (2003) discuss the validity of some of these approaches.

measure to form homogenous clusters. Specifically, the procedure aims at segmenting the data in such a way that the within-cluster variation is minimized. Consequently, we do not need to decide on a distance measure in the first step of the analysis.

The clustering process starts by randomly assigning objects to a number of clusters.¹⁰ The objects are then successively reassigned to other clusters to minimize the within-cluster variation, which is basically the (squared) distance from each observation to the center of the associated cluster. If the reallocation of an object to another cluster decreases the within-cluster variation, this object is reassigned to that cluster.

With the hierarchical methods, an object remains in a cluster once it is assigned to it, but with k-means, cluster affiliations can change in the course of the clustering process. Consequently, k-means does not build a hierarchy as described before (Fig. 9.3), which is why the approach is also frequently labeled as non-hierarchical.

For a better understanding of the approach, let's take a look at how it works in practice. Figs. 9.10–9.13 illustrate the k-means clustering process.

Prior to analysis, we have to decide on the number of clusters. Our client could, for example, tell us how many segments are needed, or we may know from previous research what to look for. Based on this information, the algorithm randomly selects a center for each cluster (step 1). In our example, two cluster centers are randomly initiated, which CC1 (first cluster) and CC2 (second cluster) in Fig. 9.10




Fig. 9.10 k-means procedure (step 1)

¹⁰Note this holds for the algorithm's original design. SPSS does not choose centers randomly.




Fig. 9.11 k-means procedure (step 2)




Fig. 9.12 k-means procedure (step 3)




Fig. 9.13 k-means procedure (step 4)

represent.¹¹ After this (step 2), Euclidean distances are computed from the cluster centers to every single object. Each object is then assigned to the cluster center with the shortest distance to it. In our example (Fig. 9.11), objects A, B, and C are assigned to the first cluster, whereas objects D, E, F, and G are assigned to the second. We now have our initial partitioning of the objects into two clusters.

Based on this initial partition, each cluster's geometric center (i.e., its centroid) is computed (third step). This is done by computing the mean values of the objects contained in the cluster (e.g., A, B, C in the first cluster) regarding each of the variables (price consciousness and brand loyalty). As we can see in Fig. 9.12, both clusters' centers now shift into new positions (CC1' for the first and CC2' for the second cluster).

In the fourth step, the distances from each object to the newly located cluster centers are computed and objects are again assigned to a certain cluster on the basis of their minimum distance to other cluster centers (CC1' and CC2'). Since the cluster centers' position changed with respect to the initial situation in the first step, this could lead to a different cluster solution. This is also true of our example, as object E is now – unlike in the initial partition – closer to the first cluster center (CC1') than to the second (CC2'). Consequently, this object is now assigned to the first cluster (Fig. 9.13). The k-means procedure now repeats the third step and re-computes the cluster centers of the newly formed clusters, and so on. In other

¹¹Conversely, SPSS always sets one observation as the cluster center instead of picking some random point in the dataset.

words, steps 3 and 4 are repeated until a predetermined number of iterations are reached, or convergence is achieved (i.e., there is no change in the cluster affiliations).

Generally, k-means is superior to hierarchical methods as it is less affected by outliers and the presence of irrelevant clustering variables. Furthermore, k-means can be applied to very large datasets, as the procedure is less computationally demanding than hierarchical methods. In fact, we suggest definitely using k-means for sample sizes above 500, especially if many clustering variables are used. From a strictly statistical viewpoint, k-means should only be used on interval or ratio-scaled data as the procedure relies on Euclidean distances. However, the procedure is routinely used on ordinal data as well, even though there might be some distortions.


One problem associated with the application of k-means relates to the fact that the researcher has to pre-specify the number of clusters to retain from the data. This makes k-means less attractive to some and still hinders its routine application in practice. However, the VRC discussed above can likewise be used for k-means clustering (an application of this index can be found in the [Web Appendix](#) → Chap. 9). Another workaround that many market researchers routinely use is to apply a hierarchical procedure to determine the number of clusters and k-means afterwards.¹² This also enables the user to find starting values for the initial cluster centers to handle a second problem, which relates to the procedure's sensitivity to the initial classification (we will follow this approach in the example application).

Two-Step Clustering

We have already discussed the issue of analyzing mixed variables measured on different scale levels in this chapter. The *two-step cluster analysis* developed by Chiu et al. (2001) has been specifically designed to handle this problem. Like k-means, the procedure can also effectively cope with very large datasets.

The name two-step clustering is already an indication that the algorithm is based on a two-stage approach: In the first stage, the algorithm undertakes a procedure that is very similar to the k-means algorithm. Based on these results, the two-step procedure conducts a modified hierarchical agglomerative clustering procedure that combines the objects sequentially to form homogenous clusters. This is done by building a so-called cluster feature tree whose “leaves” represent distinct objects in the dataset. The procedure can handle categorical and continuous variables simultaneously and offers the user the flexibility to specify the cluster numbers as well as the maximum number of clusters, or to allow the technique to automatically choose the number of clusters on the basis of statistical evaluation criteria. Likewise, the procedure guides the decision of how many clusters to retain from the data by calculating measures-of-fit such as *Akaike's Information Criterion (AIC)* or *Bayes*

¹²See Punji and Stewart (1983) for additional information on this sequential approach.

Information Criterion (BIC). Furthermore, the procedure indicates each variable's importance for the construction of a specific cluster. These desirable features make the somewhat less popular two-step clustering a viable alternative to the traditional methods. You can find a more detailed discussion of the two-step clustering procedure in the  Web Appendix (\rightarrow Chap. 9), but we will also apply this method in the subsequent example.

Validate and Interpret the Cluster Solution

Before interpreting the cluster solution, we have to assess the solution's stability and validity. Stability is evaluated by using different clustering procedures on the same data and testing whether these yield the same results. In hierarchical clustering, you can likewise use different distance measures. However, please note that it is common for results to change even when your solution is adequate. How much variation you should allow before questioning the stability of your solution is a matter of taste. Another common approach is to split the dataset into two halves and to thereafter analyze the two subsets separately using the same parameter settings. You then compare the two solutions' cluster centroids. If these do not differ significantly, you can presume that the overall solution has a high degree of stability. When using hierarchical clustering, it is also worthwhile changing the order of the objects in your dataset and re-running the analysis to check the results' stability. The results should not, of course, depend on the order of the dataset. If they do, you should try to ascertain if any obvious outliers may influence the results of the change in order.

Assessing the solution's reliability is closely related to the above, as reliability refers to the degree to which the solution is stable over time. If segments quickly change their composition, or its members their behavior, targeting strategies are likely not to succeed. Therefore, a certain degree of stability is necessary to ensure that marketing strategies can be implemented and produce adequate results. This can be evaluated by critically revisiting and replicating the clustering results at a later point in time.

To validate the clustering solution, we need to assess its criterion validity. In research, we could focus on criterion variables that have a theoretically based relationship with the clustering variables, but were not included in the analysis. In market research, criterion variables usually relate to managerial outcomes such as the sales per person, or satisfaction. If these criterion variables differ significantly, we can conclude that the clusters are distinct groups with criterion validity.

To judge validity, you should also assess face validity and, if possible, expert validity. While we primarily consider criterion validity when choosing clustering variables, as well as in this final step of the analysis procedure, the assessment of face validity is a process rather than a single event. The key to successful segmentation is to critically revisit the results of different cluster analysis set-ups (e.g., by using

different algorithms on the same data) in terms of managerial relevance. This underlines the exploratory character of the method. The following criteria will help you make an evaluation choice for a clustering solution (Dibb 1999; Tonks 2009; Kotler and Keller 2009).

- *Substantial*: The segments are large and profitable enough to serve.
- *Accessible*: The segments can be effectively reached and served, which requires them to be characterized by means of observable variables.
- *Differentiable*: The segments can be distinguished conceptually and respond differently to different marketing-mix elements and programs.
- *Actionable*: Effective programs can be formulated to attract and serve the segments.
- *Stable*: Only segments that are stable over time can provide the necessary grounds for a successful marketing strategy.
- *Parsimonious*: To be managerially meaningful, only a small set of substantial clusters should be identified.
- *Familiar*: To ensure management acceptance, the segments composition should be comprehensible.
- *Relevant*: Segments should be relevant in respect of the company's competencies and objectives.
- *Compactness*: Segments exhibit a high degree of within-segment homogeneity and between-segment heterogeneity.
- *Compatibility*: Segmentation results meet other managerial functions' requirements.

The final step of any cluster analysis is the interpretation of the clusters.

Interpreting clusters always involves examining the cluster centroids, which are the clustering variables' average values of all objects in a certain cluster. This step is of the utmost importance, as the analysis sheds light on whether the segments are conceptually distinguishable. Only if certain clusters exhibit significantly different means in these variables are they distinguishable – from a data perspective, at least. This can easily be ascertained by comparing the clusters with independent t-tests samples or ANOVA (see Chap. 6).

By using this information, we can also try to come up with a meaningful name or label for each cluster; that is, one which adequately reflects the objects in the cluster. This is usually a very challenging task. Furthermore, clustering variables are frequently unobservable, which poses another problem. How can we decide to which segment a new object should be assigned if its unobservable characteristics, such as personality traits, personal values or lifestyles, are unknown? We could obviously try to survey these attributes and make a decision based on the clustering variables. However, this will not be feasible in most situations and researchers therefore try to identify observable variables that best mirror the partition of the objects. If it is possible to identify, for example, demographic variables leading to a very similar partition as that obtained through the segmentation, then it is easy to assign a new object to a certain segment on the basis of these demographic

characteristics. These variables can then also be used to characterize specific segments, an action commonly called *profiling*.

For example, imagine that we used a set of items to assess the respondents' values and learned that a certain segment comprises respondents who appreciate self-fulfillment, enjoyment of life, and a sense of accomplishment, whereas this is not the case in another segment. If we were able to identify explanatory variables such as gender or age, which adequately distinguish these segments, then we could partition a new person based on the modalities of these observable variables whose traits may still be unknown.

Table 9.11 summarizes the steps involved in a hierarchical and k-means clustering.

While companies often develop their own market segments, they frequently use standardized segments, which are based on established buying trends, habits, and customers' needs and have been specifically designed for use by many products in mature markets. One of the most popular approaches is the PRIZM lifestyle segmentation system developed by Claritas Inc., a leading market research company. PRIZM defines every US household in terms of 66 demographically and behaviorally distinct segments to help marketers discern those consumers' likes, dislikes, lifestyles, and purchase behaviors.

Visit the Claritas website and flip through the various segment profiles. By entering a 5-digit US ZIP code, you can also find a specific neighborhood's top five lifestyle groups.

One example of a segment is "Gray Power," containing middle-class, home-owning suburbanites who are aging in place rather than moving to retirement communities. Gray Power reflects this trend, a segment of older, midscale singles and couples who live in quiet comfort.



<http://www.claritas.com/MyBestSegments/Default.jsp>

We also introduce steps related to two-step clustering which we will further introduce in the subsequent example.

Table 9.11 Steps involved in carrying out a factor analysis in SPSS

Theory	Action
<i>Research problem</i>	
Identification of homogenous groups of objects in a population	
Select clustering variables that should be used to form segments	Select relevant variables that potentially exhibit high degrees of criterion validity with regard to a specific managerial objective.
<i>Requirements</i>	
Sufficient sample size	Make sure that the relationship between objects and clustering variables is reasonable (rough guideline: number of observations should be at least 2^m , where m is the number of clustering variables). Ensure that the sample size is large enough to guarantee substantial segments.
Low levels of collinearity among the variables	► Analyze ► Correlate ► Bivariate Eliminate or replace highly correlated variables (correlation coefficients > 0.90).
<i>Specification</i>	
Choose the clustering procedure	If there is a limited number of objects in your dataset or you do not know the number of clusters: ► Analyze ► Classify ► Hierarchical Cluster If there are many observations (> 500) in your dataset and you have a priori knowledge regarding the number of clusters: ► Analyze ► Classify ► K-Means Cluster If there are many observations in your dataset and the clustering variables are measured on different scale levels: ► Analyze ► Classify ► Two-Step Cluster
Select a measure of similarity or dissimilarity (only hierarchical and two-step clustering)	<i>Hierarchical methods:</i> ► Analyze ► Classify ► Hierarchical Cluster ► Method ► Measure Depending on the scale level, select the measure; convert variables with multiple categories into a set of binary variables and use matching coefficients; standardize variables if necessary (on a range of 0 to 1 or -1 to 1). <i>Two-step clustering:</i> ► Analyze ► Classify ► Two-Step Cluster ► Distance Measure Use Euclidean distances when all variables are continuous; for mixed variables, use log-likelihood.
Choose clustering algorithm (only hierarchical clustering)	► Analyze ► Classify ► Hierarchical Cluster ► Method ► Cluster Method Use Ward's method if equally sized clusters are expected and no outliers are present. Preferably use single linkage, also to detect outliers.
Decide on the number of clusters	<i>Hierarchical clustering:</i> Examine the dendrogram: ► Analyze ► Classify ► Hierarchical Cluster ► Plots ► Dendrogram

(continued)

Table 9.11 (continued)

Theory	Action
	Draw a scree plot (e.g., using Microsoft Excel) based on the coefficients in the agglomeration schedule.
	Compute the VRC using the ANOVA procedure: ► Analyze ► Compare Means ► One-Way ANOVA
	Move the cluster membership variable in the Factor box and the clustering variables in the Dependent List box.
	Compute VRC for each segment solution and compare values.
	<i>k-means:</i> Run a hierarchical cluster analysis and decide on the number of segments based on a dendrogram or scree plot; use this information to run k-means with k clusters.
	Compute the VRC using the ANOVA procedure: ► Analyze ► Classify ► K-Means Cluster ► Options ► ANOVA table; Compute VRC for each segment solution and compare values.
	<i>Two-step clustering:</i> Specify the maximum number of clusters: ► Analyze ► Classify ► Two-Step Cluster ► Number of Clusters
	Run separate analyses using AIC and, alternatively, BIC as clustering criterion: ► Analyze ► Classify ► Two-Step Cluster ► Clustering Criterion
<i>Validate and interpret the cluster solution</i>	Examine the auto-clustering output.
Assess the solution's stability	Re-run the analysis using different clustering procedures, algorithms or distance measures. Split the datasets into two halves and compute the clustering variables' centroids; compare centroids for significant differences (e.g., independent-samples t-test or one-way ANOVA).
Assess the solution's reliability	Change the ordering of objects in the dataset (hierarchical clustering only). Replicate the analysis using a separate, newly collected dataset.
Assess the solution's validity	<i>Criterion validity:</i> Evaluate whether there are significant differences between the segments with regard to one or more criterion variables. <i>Face and expert validity:</i> Segments should be substantial, accessible, differentiable, actionable, stable, parsimonious, familiar and relevant. Segments should exhibit high degrees of within-segment homogeneity and between-segment heterogeneity. The segmentation results should meet the requirements of other managerial functions.

(continued)

Table 9.11 (continued)

Theory	Action
Interpret the cluster solution	Examine cluster centroids and assess whether these differ significantly from each other (e.g., by means of t-tests or ANOVA; see Chap. 6). Identify names or labels for each cluster and characterize each cluster by means of observable variables, if necessary (cluster profiling).

Example

Silicon-Valley-based Tesla Motors Inc. (<http://www.teslamotors.com>) is an automobile startup company focusing on the production of high performance electrical vehicles with exceptional design, technology, performance, and efficiency. Having reported the 500th delivery of its roadster in June 2009, the company decided to focus more strongly on the European market. However, as the company has only limited experience in the European market, which has different characteristics than that of the US, it asked a market research firm to provide a segmentation concept. Consequently, the market research firm gathered data from major car manufacturers on the following car characteristics, all of which have been measured on a ratio scale (variable names in parentheses):

- Engine displacement (*displacement*)
- Turning moment in Nm (*moment*)
- Horsepower (*horsepower*)
- Length in mm (*length*)
- Width in mm (*width*)
- Net weight in kg (*weight*)
- Trunk volume in liters (*trunk*)
- Maximum speed in km/h (*speed*)
- Acceleration 0–100 km/h in seconds (*acceleration*)



The pretest sample of 15, randomly taken, cars is shown in Fig. 9.14. In practice, clustering is done on much larger samples but we use a small sample size to illustrate the clustering process. Keep in mind that in this example, the ratio between the objects and clustering variables is much too small. The dataset used is *cars.sav* (📄 Web Appendix → Chap. 9).

	Name	displacement	moment	horsepower	length	width	weight	trunk	speed	acceleration
1	Kia Picanto 1.1 Start	1086	97	65	3635	1595	929	127	154	15.10
2	Suzuki Splash 1.0	996	90	65	3715	1680	1050	178	160	14.70
3	Renault Clio 1.2	1149	105	75	3966	1719	1155	288	167	13.40
4	Dacia Sandero 1.6	1598	128	87	4020	1746	1111	320	174	11.50
5	Fiat Grande Punto 1.4	1598	140	88	3986	1719	1215	288	177	11.90
6	Peugot 207 1.4	1360	133	88	4030	1748	1214	270	180	12.70
7	Renault Clio 1.6	1368	125	95	4030	1687	1135	275	178	11.40
8	Porsche Cayman	3386	340	295	4341	1801	1340	410	275	5.40
9	Nissan 350Z	3498	353	301	4315	1815	1610	235	250	5.80
10	Mercedes C 200 CDI	2148	270	136	4595	1770	1605	495	208	10.80
11	VWPassat Variant 2.0	1968	320	140	4774	1820	1596	588	201	10.50
12	Skoda Octavia 2.0	1968	320	140	4572	1769	1425	580	207	9.70
13	Mercedes E 280	2996	300	231	4852	1822	1660	540	250	7.30
14	Audi A6 2.4	2393	230	177	4916	1855	1525	546	231	8.90
15	BMW 525i	2497	250	218	4841	1846	1550	520	245	7.50

Fig. 9.14 Data

In the next step, we will run several different clustering procedures on the basis of these nine variables. We first apply a hierarchical cluster analysis based on Euclidean distances, using the single linkage method. This will help us determine a suitable number of segments, which we will use as input for a subsequent k-means clustering. Finally, we will run a two-step cluster analysis using SPSS.

Before we start with the clustering process, we have to examine the variables for substantial collinearity. Just by looking at the variable set, we suspect that there are some highly correlated variables in our dataset. For example, we expect rather high correlations between *speed* and *acceleration*. To determine this, we run a bivariate correlation analysis by clicking ▶ Analyze ▶ Correlate ▶ Bivariate, which will open a dialog box similar to that in Fig. 9.15. Enter all variables into the **Variables** box and select the box **Pearson** (under **Correlation Coefficients**) because these are continuous variables.

The correlation matrix in Table 9.12 supports our expectations – there are several variables that have high correlations. *Displacement* exhibits high (absolute) correlation coefficients with *horsepower*, *speed*, and *acceleration*, with values well above 0.90, indicating possible collinearity issues. Similarly, *horsepower* is highly correlated with *speed* and *acceleration*. Likewise, *length* shows a high degree of correlation with *width*, *weight*, and *trunk*.

A potential solution to this problem would be to run a factor analysis and perform a cluster analysis on the resulting factor scores. Since the factors obtained are, by definition, independent, this would allow for an effective handling of the collinearity issue. However, as this approach is associated with several problems (see discussion above) and as there are only several variables in our data set, we should reduce the variables, for example, by omitting *displacement*, *horsepower*,




Fig. 9.15 Bivariate correlations dialog box

and *length* from the subsequent analyses. The remaining variables still provide a sound basis for carrying out cluster analysis.

To run the hierarchical clustering procedure, click on ► Analyze ► Classify ► Hierarchical Cluster, which opens a dialog box similar to Fig. 9.16.

Move the variables *moment*, *width*, *weight*, *trunk*, *speed*, and *acceleration* into the **Variable(s)** box and specify *name* as the labeling variable (box **Label Cases by**). The **Statistics** option gives us the opportunity to request the distance matrix (labeled proximity matrix in this case) and the agglomeration schedule, which provides information on the objects being combined at each stage of the clustering process. Furthermore, we can specify the number or range of clusters to retain from the data. As we do not yet know how many clusters to retain, just check the box **Agglomeration schedule** and continue.

Under **Plots**, we choose to display a dendrogram, which graphically displays the distances at which objects and clusters are joined. Also ensure you select the icicle diagram (for all clusters), which is yet another graph for displaying clustering solutions.

The option **Method** allows us to specify the cluster method (e.g., single linkage or Ward's method), the distance measure (e.g., Chebychev distance or the Jaccard coefficient), and the type of standardization of values. In this example, we use the single linkage method (**Nearest neighbor**) based on **Euclidean distances**. Since the variables are measured on different levels (e.g., speed versus weight), make sure to standardize the variables, using, for example, the **Range –1 to 1 (by variable)** in the **Transform Values** drop-down list.

Table 9.12 Correlation matrix

Correlations										
	displacement	moment	horsepower	length	width	weight	trunk	speed	acceleration	
displacement										
Pearson Correlation	1	.875**	.983**	.657**	.764**	.768**	.470	.967**	-.969**	
Sig. (2-tailed)		.000	.000	.008	.001	.001	.077	.000	.000	
N	15	15	15	15	15	15	15	15	15	15
moment										
Pearson Correlation	.875**	1	.847**	.767**	.766**	.862**	.691**	.859**	-.861**	
Sig. (2-tailed)	.000		.000	.001	.001	.000	.004	.000	.000	
N	15	15	15	15	15	15	15	15	15	15
horsepower										
Pearson Correlation	.983**	.847**	1	.608*	.732**	.714**	.408	.968**	-.961**	
Sig. (2-tailed)	.000	.000		.016	.002	.003	.131	.000	.000	
N	15	15	15	15	15	15	15	15	15	15
length										
Pearson Correlation	.657**	.767**	.608*	1	.912**	.921**	.934**	.741**	-.714**	
Sig. (2-tailed)	.008	.001	.016		.000	.000	.000	.002	.003	
N	15	15	15	15	15	15	15	15	15	15
width										
Pearson Correlation	.764**	.766**	.732**	.912**	1	.884**	.783**	.819**	-.818**	
Sig. (2-tailed)	.001	.001	.002	.000		.000	.001	.000	.000	
N	15	15	15	15	15	15	15	15	15	15
weight										
Pearson Correlation	.768**	.862**	.714**	.921**	.884**	1	.785**	.778**	-.763**	
Sig. (2-tailed)	.001	.000	.003	.000	.000		.001	.001	.001	
N	15	15	15	15	15	15	15	15	15	15
trunk										
Pearson Correlation	.470	.691**	.408	.934**	.783**	.785**	1	.579*	-.552*	
Sig. (2-tailed)	.077	.004	.131	.000	.001	.001		.024	.033	
N	15	15	15	15	15	15	15	15	15	15
speed										
Pearson Correlation	.967**	.859**	.968**	.741**	.819**	.778**	.579*	1	-.971**	
Sig. (2-tailed)	.000	.000	.000	.002	.000	.001	.024		.000	
N	15	15	15	15	15	15	15	15	15	15
acceleration										
Pearson Correlation	-.969**	-.861**	-.961**	-.714**	-.818**	-.763**	-.552*	-.971**	1	
Sig. (2-tailed)	.000	.000	.000	.003	.000	.001	.033	.000	.000	
N	15	15	15	15	15	15	15	15	15	15

** . Correlation is significant at the 0.01 level (2-tailed).

* . Correlation is significant at the 0.05 level (2-tailed).

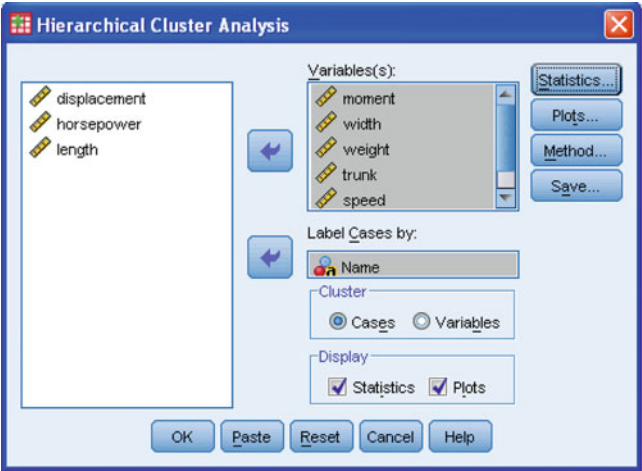


Fig. 9.16 Hierarchical cluster analysis dialog box

Lastly, the **Save** option enables us to save cluster memberships for a single solution or a range of solutions. Saved variables can then be used in subsequent analyses to explore differences between groups. As a start, we will skip this option, so continue and click on **OK** in the main menu.

Table 9.13 Agglomeration schedule

Agglomeration Schedule						
Stage	Cluster Combined		Coefficients	Stage Cluster First Appears		Next Stage
	Cluster 1	Cluster 2		Cluster 1	Cluster 2	
1	5	6	.149	0	0	2
2	5	7	.184	1	0	3
3	4	5	.201	0	2	5
4	14	15	.213	0	0	6
5	3	4	.220	0	3	8
6	13	14	.267	0	4	11
7	11	12	.321	0	0	9
8	2	3	.353	0	5	10
9	10	11	.357	0	7	11
10	1	2	.389	0	8	14
11	10	13	.484	9	6	13
12	8	9	.575	0	0	13
13	8	10	.618	12	11	14
14	1	8	.910	10	13	0

First, we take a closer look at the agglomeration schedule (Table 9.13), which displays the objects or clusters combined at each stage (second and third column) and the distances at which this merger takes place. For example, in the first stage, objects 5 and 6 are merged at a distance of 0.149. From here onward, the resulting cluster is labeled as indicated by the first object involved in this merger, which is object 5. The last column on the very right tells you in which stage of the algorithm this cluster will appear next. In this case, this happens in the second step, where it is merged with object 7 at a distance of 0.184. The resulting cluster is still labeled 5, and so on. Similar information is provided by the icicle diagram shown in Fig. 9.17. Its name stems from the analogy to rows of icicles hanging from the eaves of a house. The diagram is read from the bottom to the top, therefore the columns correspond to the objects being clustered, and the rows represent the number of clusters.




Fig. 9.17 Icicle diagram

As described earlier, we can use the agglomeration schedule to determine the number of segments to retain from the data. By plotting the distances (**Coefficients** column in Table 9.13) against the number of clusters, using a spreadsheet program, we can generate a scree plot. The distinct break (elbow) generally indicates the solution regarding where an additional combination of two objects or clusters would occur at a greatly increased distance. Thus, the number of clusters prior to this merger is the most probable solution. The scree plot - which we made



Fig. 9.18 Scree plot

separately using Microsoft Excel - (Fig. 9.18) does not show such a distinct break. Note that – unlike in the factor analysis – we do not pick the solution with one cluster less than indicated by the elbow. The sharp increase in distance when switching from a one to a two-cluster solution occurs in almost all analyses and must not be viewed as a reliable indicator for the decision regarding the number of segments.

The scree plot in Fig. 9.18 shows that there is no clear elbow indicating a suitable number of clusters to retain. Based on the results, one could argue for a five-segment or six-segment solution. However, considering that there are merely 15 objects in the dataset, this seems too many, as we then have very small (and, most probably, meaningless) clusters. Consequently, a two, three or four-segment solution is deemed more appropriate.

Let’s take a look at the dendrogram shown in Fig. 9.19. We read the dendrogram from left to right. Vertical lines are objects and clusters joined together – their position indicates the distance at which this merger takes place. When creating a dendrogram, SPSS rescales the distances to a range of 0–25; that is, the last merging step to a one-cluster solution takes place at a (rescaled) distance of 25. Note that this differs from our manual calculation shown in Fig. 9.9, where we did not do any rescaling! Again, the analysis only provides a rough guidance regarding the number of segments to retain. The change in distances between the mergers indicates that (besides a two-segment solution) both a three and four-segment solution are appropriate.

To clarify this issue, let’s re-run the analysis, but this time we pre-specify different segment numbers to compare these with regard to content validity. To do so, just re-run the analysis using hierarchical clustering. Now switch to the **Save** option, specify a range of solutions from 2 to 4 and run the analysis. SPSS generates




Fig. 9.19 Dendrogram

the same output but also adds three additional variables to your dataset (*CLU4_I*, *CLU3_I*, and *CLU2_I*), which reflect each object’s cluster membership for the respective analysis. SPSS automatically places *CLU* in front, followed by the number of clusters (4, 3, or 2), to identify each object’s cluster membership. The results are

Table 9.14 Cluster memberships

Name	Four clusters, observation member of cluster	Three clusters, observation member of cluster	Two clusters, observation member of cluster
Kia Picanto 1.1 Start	1	1	1
Suzuki Splash 1.0	1	1	1
Renault Clio 1.2	1	1	1
Dacia Sandero 1.6	1	1	1
Fiat Grande Punto 1.4	1	1	1
Peugot 207 1.4	1	1	1
Renault Clio 1.6	1	1	1
Porsche Cayman	2	2	2
Nissan 350Z	3	2	2
Mercedes C 200 CDI	4	3	2
VW Passat Variant 2.0	4	3	2
Skoda Octavia 2.0	4	3	2
Mercedes E 280	4	3	2
Audi A6 2.4	4	3	2
BMW 525i	4	3	2

illustrated in Table 9.14. SPSS does not produce this table for us, so we need to enter these cluster memberships ourselves in a table or spreadsheet.

When we view the results, a three-segment solution appears promising. In this solution, the first segment comprises compact cars, whereas the second segment contains sports cars, and the third limousines. Increasing the solution by one segment would further split up the sports cars segment into two sub-segments. This does not appear to be very helpful, as now two of the four segments comprise only one object. This underlines the single linkage method’s tendency to identify outlier objects – in this case the Nissan 350Z and Porsche Cayman. In this specific example, the Nissan 350Z and Porsche Cayman should not be regarded as outliers in a classical sense but rather as those cars which may be Tesla’s main competitors in the sports car market.

In contrast, the two-segment solution appears to be rather imprecise considering the vast differences in the mix of sports and middle-sized cars in this solution.

To get a better overview of the results, let’s examine the cluster centroids; that is, the mean values of the objects contained in the cluster on selected variables. To do so, we split up the dataset using the **Split File** command (► Data ► Split File) (see Chap. 5). This enables us to analyze the data on the basis of a grouping variable’s values. In this case, we choose *CLU3_1* as the grouping variable and select the option **Compare groups**. Subsequently, we calculate descriptive statistics (► Analyze ► Descriptive Statistics ► Descriptives, also see Chap. 5) and calculate the mean, minimum and maximum values, as well as the standard deviations of the clustering variables. Table 9.15 shows the results for the variables *weight*, *speed*, and *acceleration*.

Table 9.15 Cluster centroids

Descriptive Statistics						
CLU3_1		N	Minimum	Maximum	Mean	Std. Deviation
1	weight	7	929	1215	1115.57	100.528
	speed	7	154	180	170.00	9.950
	acceleration	7	11.40	15.10	12.9571	1.50317
	Valid N (listwise)	7				
2	weight	2	1340	1610	1475.00	190.919
	speed	2	250	275	262.50	17.678
	acceleration	2	5.40	5.80	5.6000	.28284
	Valid N (listwise)	2				
3	weight	6	1425	1660	1560.17	81.081
	speed	6	201	250	223.67	21.163
	acceleration	6	7.30	10.80	9.1167	1.48649
	Valid N (listwise)	6				

From the descriptive statistics, it seems that the first segment contains light-weight compact cars (with a lower maximum speed and acceleration). In contrast, the second segment comprises two sports cars with greater speed and acceleration, whereas the third contains limousines with an increased weight and intermediate speed and acceleration. Since the descriptives do not tell us if these differences are significant, we could have used a one-way ANOVA (menu ► Analyze ► Compare Means ► One-Way ANOVA) to calculate the cluster centroids and compare the differences formally.

In the next step, we want to use the *k-means* method on the data. We have previously seen that we need to specify the number of segments when conducting k-means clustering. SPSS then initiates cluster centers and assigns objects to the clusters based on their minimum distance to these centers. Instead of letting SPSS choose the centers, we can also save the centroids (cluster centers) from our previous analysis as input for the k-means procedure. To do this, we need to do some data management in SPSS, as the cluster centers have to be supplied in a specific format. Consequently, we need to aggregate the data first (briefly introduced in Chap. 5). By selecting ► Data ► Aggregate, a dialog box similar to Fig. 9.20 opens up.




Fig. 9.20 Aggregate data dialog box

Note that we choose **Display Variable Names** instead of **Display Variable Labels** by clicking the right mouse button on the left box showing the variables in the dataset. Now we proceed by choosing the cluster membership variable (*CLU3_1*) as a break variable and move the *moment*, *width*, *weight*, *trunk*, *speed*, and *acceleration* variables into the **Summaries of Variable(s)** box. When using the default settings, SPSS computes the variables' mean values along the lines of the break variable (indicated by the postfix *_mean*, which is added to each aggregate variable's name), which corresponds to the cluster centers that we need for the k-means analysis. You can change each aggregate variable's name from the original one by removing the postfix *_mean* – using the **Name & Label** option – if you want to. Lastly, we do not want to add the aggregated variables to the active dataset, but rather need to create a new dataset comprising only the aggregated variables. You must therefore check this under **SAVE** and specify a dataset label such as *aggregate*. When clicking on **OK**, a new dataset labeled *aggregate* is created and opened automatically.

The new dataset is almost in the right format – but we still need to change the break variable's name from *CLU3_1* to *cluster_* (SPSS will issue a warning but this can be safely ignored). The final dataset should have the form shown in Fig. 9.21.

Now let's proceed by using k-means clustering. Make sure that you open the original dataset and go to Analyze ► Classify ► K-Means Cluster, which brings up a new dialog box (Fig. 9.22).

	cluster_	moment	width	weight	trunk	speed	acceleration
1	1	116.86	1699.14	1115.57	249.43	170.00	12.96
2	2	346.50	1808.00	1475.00	322.50	262.50	5.60
3	3	281.67	1813.67	1560.17	543.17	223.67	9.12

Fig. 9.21 Aggregated data file

As you did in the hierarchical clustering analysis, move the six clustering variables to the **Variables** box and specify the case labels (variable name). To use the cluster centers from our previous analysis, check the box **Read initial** and click on **Open dataset**. You can now choose the dataset labeled *aggregate*. Specify 3, which corresponds to the result of the hierarchical clustering analysis, in the **Number of Clusters** box. The **Iterate** option is of less interest to us. Instead, click on **Save** and check the box **Cluster Membership**. This creates a new variable indicating each object's final cluster membership. SPSS indicates whether each observation is a member of cluster 1, 2, or 3. Under **Options**, you can request several statistics and specify how missing values should be treated. Ensure that you request the initial cluster centers as well as the ANOVA table and that you exclude the missing values listwise (default). Now start the analysis.

The k-means procedure generates Tables 9.16 and 9.17, which show the initial and final cluster centers. As you can see, these are identical (also compare Fig. 9.21), which indicates that the initial partitioning of the objects in the first step of the k-means procedure was retained during the analysis. This means that it

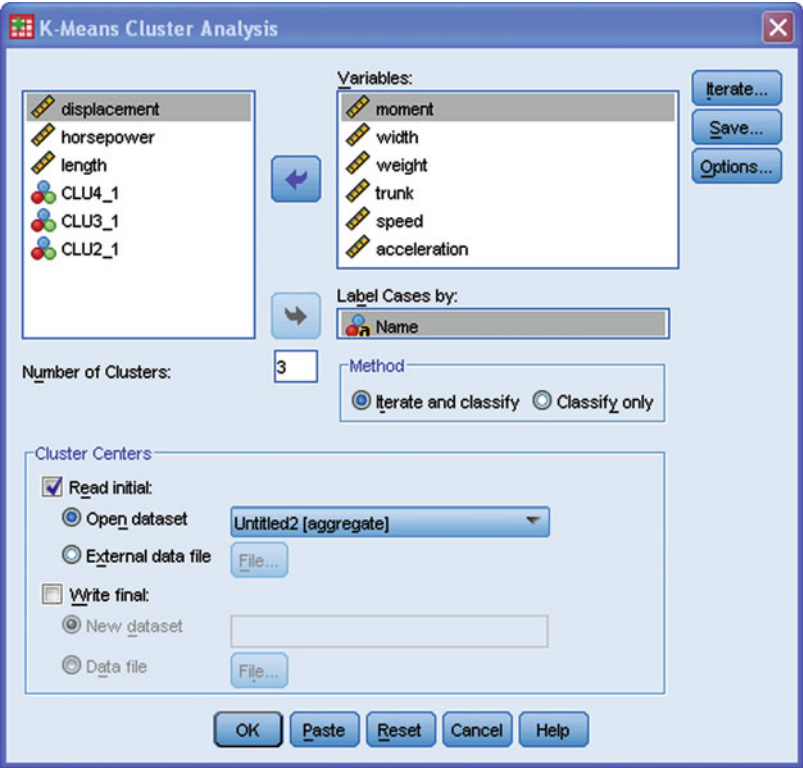


Fig. 9.22 K-means cluster analysis dialog box

Table 9.16 Initial cluster centers

	Initial Cluster Centers		
	Cluster		
	1	2	3
moment	117	347	282
width	1699	1808	1814
weight	1116	1475	1560
trunk	249	323	543
speed	170	263	224
acceleration	12.96	5.60	9.12

Input from FILE Subcommand

Table 9.17 Final cluster centers

Final Cluster Centers			
	Cluster		
	1	2	3
moment	117	347	282
width	1699	1808	1814
weight	1116	1475	1560
trunk	249	323	543
speed	170	263	224
acceleration	12.96	5.60	9.12

was not possible to reduce the overall within-cluster variation by re-assigning objects to different clusters.


Likewise, the output **Iteration History** shows that there is no change in the cluster centers. Similarly, if you compare the partitioning of objects into the three clusters by examining the newly generated variable *QCL_I*, you see that there is no change in the clusters’ composition. At first sight, this does not look like a very exciting result, but this in fact signals that the initial clustering solution is stable. In other words, the fact that two different clustering methods yield the same outcomes provides some evidence of the results’ stability.

In contrast to hierarchical clustering, the k-means outputs provide us with an ANOVA of the cluster centers (Table 9.18). As you can see, all the clustering variables’ means differ significantly across at least two of the three segments, because the null hypothesis is rejected in every case (**Sig.** ≤ 0.05).

Table 9.18 ANOVA output

	ANOVA					
	Cluster		Error		F	Sig.
	Mean Square	df	Mean Square	df		
moment	64318.455	2	784.224	12	82.015	.000
width	23904.771	2	1966.183	12	12.158	.001
weight	339920.393	2	10829.712	12	31.388	.000
trunk	142764.143	2	4311.754	12	33.110	.000
speed	8628.283	2	262.153	12	32.913	.000
acceleration	50.855	2	2.057	12	24.722	.000

Since we used the prior analysis results from hierarchical clustering as an input for the k-means procedure, the problem of selecting the “correct” number of segments is not problematic in this example. As discussed above, we could have

also used the VRC to make that decision. In the  Web Appendix (→ Chap. 9), we present a VRC application to this example.

As a last step of the analysis, we conduct a two-step clustering approach. First, go to Analyze ► Classify ► Two-Step Cluster. A new dialog box is opened, similar to that shown in Fig. 9.23.

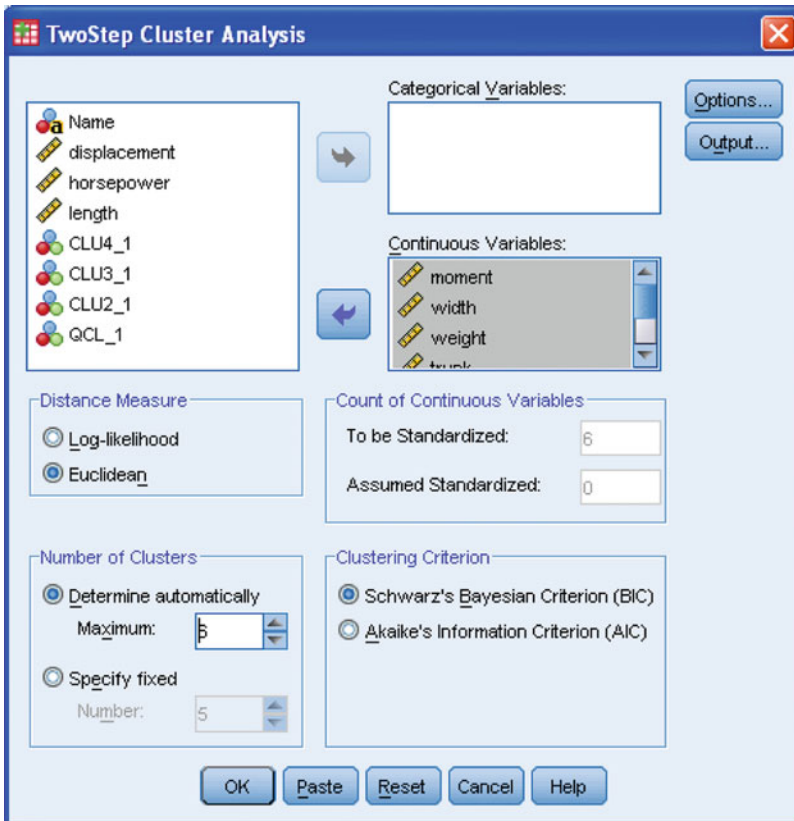



Fig. 9.23 Two-step cluster analysis dialog box

Move the variables we used in the previous analyses to the **Continuous Variables** box.

The **Distance Measure** box determines how the distance between two objects or clusters is computed. While **Log-likelihood** can be used for categorical and continuous variables, the **Euclidean** distance can only be applied when all of the variables are continuous. Unless your dataset contains categorical variables (e.g., gender) you should choose the Euclidean distance measure, as this generally provides better results. If you use ordinal variables and therefore use the **Log-likelihood** procedure, check that the answer categories are equidistant. In our dataset, all variables are continuous, therefore select the second option, namely **Euclidean**.

Under **Number of Clusters**, you can specify a fixed number or a maximum number of segments to retain from the data. One of two-step clustering's major advantages is that it allows the automatic selection of the number of clusters. To make use of this advantage, you should specify a maximum number of clusters, for example, 6. Next to this box, in which the number of clusters is specified, you can choose between two criteria (also referred to as model selection or information criteria) which SPSS can use to pick an appropriate number of segments, namely **Akaike's information criterion (AIC)** and **Bayes information criterion (BIC)**. These are relative measures of goodness-of-fit and are used to compare different solutions with different numbers of segments. "Relative" means that these criteria are not scaled on a range of, for example, 0 to 1 but can generally take any value. Compared to an alternative solution with a different number of segments, smaller values in AIC or BIC indicate an increased fit. SPSS computes solutions for different segment numbers (up to the maximum number of segments specified before) and chooses the appropriate solution by looking for the smallest value in the chosen criterion. However, which criterion should we choose? AIC is well-known for overestimating the "correct" number of segments, while BIC has a slight tendency to underestimate this number. Thus, it is worthwhile comparing the clustering outcomes of both criteria and selecting a smaller number of segments than actually indicated by AIC. Nevertheless, when running two separate analyses, one based on AIC and the other based on BIC, SPSS usually renders the same results. But what do we do if the two criteria indicate different numbers of clusters? In such a situation, we should evaluate each solution on practical grounds as well as in light of the solution's interpretability. Do not solely rely on the automatic model selection, especially when there is a combination of continuous and categorical variables, as this does not always work well. Examine the results very carefully!

Under **Options**, you can specify issues related to outlier treatment, memory allocation, and variable standardization. Variables that are already standardized have to be assigned as such, but since this is not the case in our analysis, you can simply proceed.

Lastly, under the option **Output**, we can specify additional variables for describing the obtained clusters. However, let's stick to the default option for now. Make sure that you click the box **Create cluster membership variable** before clicking **Continue**. Note that the menu options as well as the outputs in SPSS 18 are no longer the same as in prior SPSS versions. Here, we discuss the menus and outputs as provided in SPSS 18, but if you want to learn what the analysis looks like in SPSS 17 (and prior versions), go to the  Web Appendix (→ Chap. 9).

SPSS produces a very simple output, as shown in Fig. 9.24. The upper part of the output describes the algorithm applied, the number of variables used (labeled input features) and the final number of clusters retained from the data. In our case, the number of clusters is chosen according to BIC, which indicates a two-segment solution (the same holds when using AIC instead of BIC). Note that this result differs from our previous analysis!

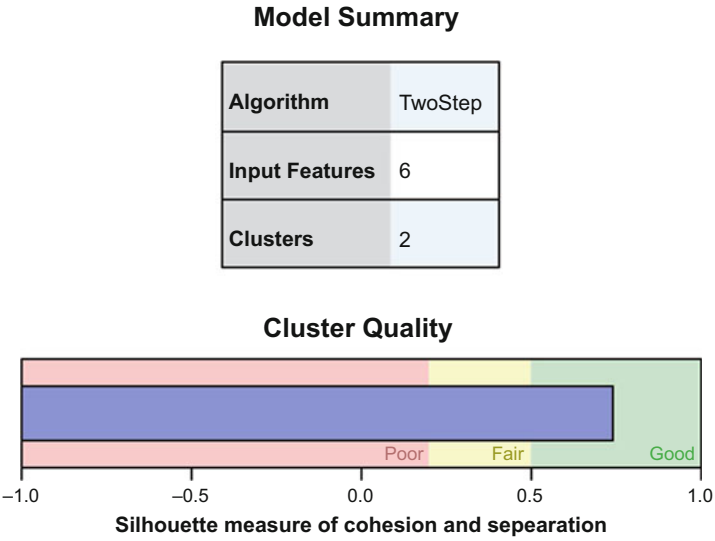


Fig. 9.24 Two-step clustering output

The lower part of the output (Fig. 9.24) indicates the quality of the cluster solution. The *silhouette measure of cohesion and separation* is a measure of the clustering solution’s overall goodness-of-fit. It is essentially based on the average distances between the objects and can vary between -1 and $+1$. Specifically, a silhouette measure of less than 0.20 indicates a poor solution quality, a measure between 0.20 and 0.50 a fair solution, whereas values of more than 0.50 indicate a good solution (this is also indicated under the horizontal bar in Fig. 9.24). In our case, the measure indicates a satisfactory cluster quality. Consequently, you can proceed with the analysis by double-clicking on the output. This will open up the model viewer (Fig. 9.25), an evaluation tool that graphically presents the structure of the revealed clusters.

The model viewer provides us with two windows: the main view, which initially shows a model summary (left-hand side), and an auxiliary view, which initially features the cluster sizes (right-hand side). At the bottom of each window, you can request different information, such as an overview of the cluster structure and the overall variable importance as shown in Fig. 9.25.

In the main view, we can see a description of the two clusters, including their (relative) sizes. Furthermore, the output shows each clustering variables’ mean values across the two clusters as well as their relative importance. Darker shades (i.e., higher values in feature importance) denote the variable’s greater importance for the clustering solution. Comparing the results, we can see that *moment* is the most important variable for each of the clusters, followed by *weight*, *speed*, *width*, *acceleration*, and *trunk*. Clicking on one of the boxes will show a graph with the frequency distribution of each cluster. The auxiliary view shows an overview

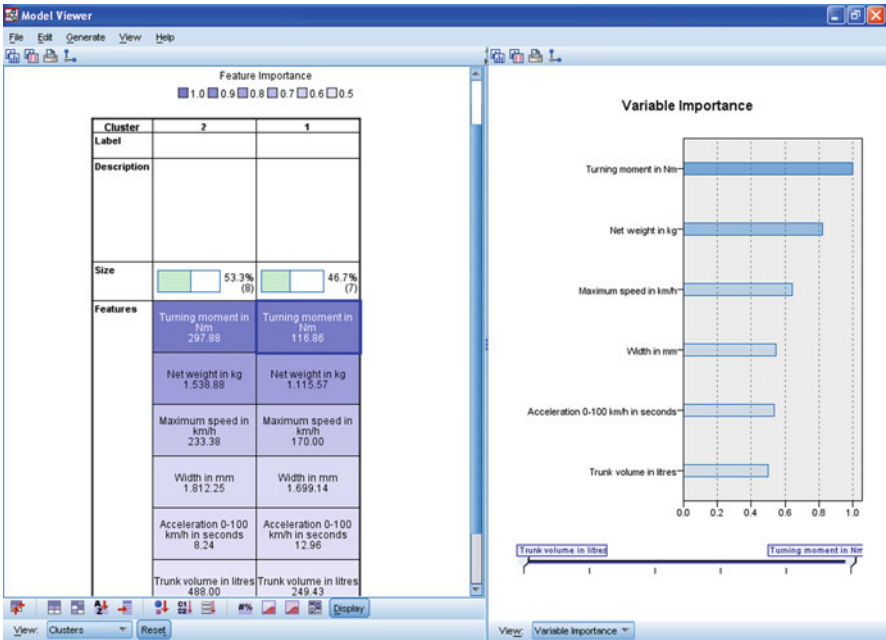


Fig. 9.25 Additional options in the model viewer

of the variables’ overall importance for the clustering solution, which provides the same result as the cluster-specific analysis. The model viewer provides us with additional options for visualizing the results or comparing clustering solutions. It is worthwhile to simply play around with the different self-explanatory options. So go ahead and explore the model viewer’s features yourself!

Case Study

Facing dramatically declining sales and decreased turnover, retailers such as Saks Fifth Avenue and JCPenney are rethinking their pricing strategies, scaling back inventories, and improving the fashion content. Men’s accessories are one of the bright spots and Saks Fifth Avenue has jumped on the trend with three recently opened shops prominently featuring this category. The largest men’s store opened in Beverly Hills in the late 2008 and stocks top brands in jewelry, watches, sunglasses, and leather goods. By providing a better showcase for men’s accessories, Saks aims at strengthening its position in a market that is often neglected in the larger department store arena. This is because the men’s accessories business generally requires expertise in buying since this typically involves small, artisan vendors – an investment many department stores are not willing to make.

The Beverly Hills store was chosen to spearhead the accessories program because it is considered the company's West Coast flagship and the department had not had a significant facelift since the store opened in 1995.¹³

Saks's strategy seemed to be successful if one considers that the newly opened boutiques already exerted an impact on sales during their first holiday season. However, before opening accessories shops in any other existing Saks stores, the company wanted to gain further insights into their customers' preferences. Consequently, a survey was conducted among visitors of the Beverly Hills store to gain a deeper understanding of their attitudes to buying and shopping. Overall, 180 respondents were interviewed using mall-intercept interviewing. The respondents were asked to indicate the importance of the following factors when buying products and services using a 5-point scale (1 = not at all important, 5 = very important):

- Saving time (x_1)
- Getting bargains (x_2)
- Getting products that aren't on the high street (x_3)
- Trying new things (x_4)
- Being aware of what companies have to offer (x_5)

The resulting dataset *Buying Attitudes.sav* (📄 Web Appendix → Chap. 9) also includes each respondent's gender and monthly disposable income.¹⁴

1. Given the levels of measurement, which clustering method would you prefer? Carry out a cluster analysis using this procedure.
2. Interpret and profile the obtained clusters by examining cluster centroids. Compare differences across clusters on observed variables using ANOVA and post-hoc tests (see Chap. 6).
3. Use a different clustering method to test the stability of your results. If necessary, omit or rescale certain variables.
4. Based on your evaluation of the dataset, make recommendations to the management of Saks's Beverly Hills store.

Questions

1. In your own words, explain the objective and basic concept of cluster analysis.
2. What are the differences between hierarchical and partitioning methods? When do we use hierarchical or partitioning methods?
3. Run the k-means analysis again from the example application (*Cars.sav*, 📄 Web Appendix → Chap. 9). Compute a three-segment solution and compare the results with those obtained by the initial hierarchical clustering.

¹³For further information, see Palmieri JE (2008). "Saks Adds Men's Accessories Shops," *Women's Wear Daily*, 196 (128), 14.

¹⁴Note that the data are artificial.

4. Run the k-means analysis again from the example application (*Cars.sav*, Web Appendix → Chap. 9). Use a factor analysis considering all nine variables and perform a cluster analysis on the resulting factor scores (factor-cluster segmentation). Interpret the results and compare these with the initial analysis.
5. Repeat the manual calculations of the hierarchical clustering procedure from the beginning of the chapter, but use complete or average linkage as clustering method. Compare the results with those of the single linkage method.
6. Make a list of the market segments to which you belong! What clustering variables did you take into consideration when you placed yourself in those segments?

Further Readings

Bottomley P, Nairn A (2004) Blinded by science: The managerial consequences of inadequately validated cluster analysis solutions. *Int J Mark Res* 46(2):171–187

In this article, the authors investigate if managers could distinguish between cluster analysis outputs derived from real-world and random data. They show that some managers feel able to assign meaning to random data devoid of a meaningful structure, and even feel confident formulating entire marketing strategies from cluster analysis solutions generated from such data. As such, the authors provide a reminder of the importance of validating clustering solutions with caution.

Everitt BS, Landau S, Leese M (2001) Cluster analysis, 4th edn. Arnold, London

This book is comprehensive yet relatively non-mathematical, focusing on the practical aspects of cluster analysis. The authors discuss classical approaches as well as more recent methods such as finite mixture modeling and neural networks.

Journal of Classification. New York, NY: Springer, available at:

<http://www.springer.com/statistics/statistical+theory+and+methods/journal/357>

If you are interested in the most recent advances in clustering techniques and have a strong background in statistics, you should check out this journal. Among the disciplines represented are statistics, psychology, biology, anthropology, archeology, astronomy, business, marketing, and linguistics.

Punj G, Stewart DW (1983) Cluster analysis in marketing research: review and suggestions for application. *J Mark Res* 20(2):134–148

In this seminal article, the authors discuss several issues in applications of cluster analysis and provide further theoretical discussion of the concepts and rules of thumb that we included in this chapter.

Romesburg C (2004) Cluster analysis for researchers. Lulu Press, Morrisville, NC

Charles Romesburg nicely illustrates the most frequently used methods of hierarchical cluster analysis for readers with limited backgrounds in mathematics and statistics.

Wedel M, Kamakura WA (2000) Market segmentation: conceptual and methodological foundations, 2nd edn. Kluwer Academic, Boston, NE

This book is a clear, readable, and interesting presentation of applied market segmentation techniques. The authors explain the theoretical concepts of recent

analysis techniques and provide sample applications. Probably the most comprehensive text in the market.

References

- Andrews RL, Currim IS (2003) Recovering and profiling the true segmentation structure in markets: an empirical investigation. *Int J Res Mark* 20(2):177–192
- Arabie P, Hubert L (1994) Cluster analysis in marketing research. In: Bagozzi RP (ed) *Advanced methods in marketing research*. Blackwell, Cambridge, pp 160–189
- Bishop CM (2006) *Pattern recognition and machine learning*. Springer, Berlin
- Calinski, T, Harabasz J (1974) A dendrite method for cluster analysis. *Commun Stat Theory Methods* 3(1):1–27
- Chiu T, Fang D, Chen J, Wang Y, Jeris C (2001) A robust and scalable clustering algorithm for mixed type attributes in large database environment. In: *Proceedings of the 7th ACM SIGKDD international conference in knowledge discovery and data mining*, Association for Computing Machinery, San Francisco, CA, pp 263–268
- Dibb S (1999) Criteria guiding segmentation implementation: reviewing the evidence. *J Strateg Mark* 7(2):107–129
- Dolnicar S (2003) Using cluster analysis for market segmentation – typical misconceptions, established methodological weaknesses and some recommendations for improvement. *Australas J Mark Res* 11(2):5–12
- Dolnicar S, Grün B (2009) Challenging “factor-cluster segmentation”. *J Travel Res* 47(1):63–71
- Dolnicar S, Lazarevski K (2009) Methodological reasons for the theory/practice divide in market segmentation. *J Mark Manage* 25(3–4):357–373
- Formann AK (1984) *Die Latent-Class-Analyse: Einführung in die Theorie und Anwendung*. Beltz, Weinheim
- Kaufman L, Rousseeuw PJ (2005) *Finding groups in data. An introduction to cluster analysis*. Wiley, Hoboken, NY
- Kohonen T (1982) Self-organized formation of topologically correct feature maps. *Biol Cybern* 43(1):59–69
- Kotler P, Keller KL (2009) *Marketing management*, 13th edn. Pearson Prentice Hall, Upper Saddle River, NJ
- Larson JS, Bradlow ET, Fader PS (2005) An exploratory look at supermarket shopping paths. *Int J Res Mark* 22(4):395–414
- McLachlan GJ, Peel D (2000) *Finite mixture models*. Wiley, New York, NY
- Milligan GW, Cooper M (1985) An examination of procedures for determining the number of clusters in a data set. *Psychometrika* 50(2):159–179
- Milligan GW, Cooper M (1988) A study of variable standardization. *J Classification* 5(2):181–204
- Moroko L, Uncles MD (2009) Employer branding and market segmentation. *J Brand Manage* 17(3):181–196
- Okazaki S (2006) What do we know about Mobile Internet Adopters? A Cluster Analysis. *Inf Manage* 43(2):127–141
- Punji G, Stewart DW (1983) Cluster analysis in marketing research: review and suggestions for application. *J Mark Res* 20(2):134–148
- Sheppard A (1996) The sequence of factor analysis and cluster analysis: differences in segmentation and dimensionality through the use of raw and factor scores,” *tourism analysis*. *Tourism Anal* 1(Inaugural Volume):49–57
- Tonks DG (2009) Validity and the design of market segments. *J Mark Manage* 25(3–4):341–356
- Wedel M, Kamakura WA (2000) *Market segmentation: conceptual and methodological foundations*, 2nd edn. Kluwer, Boston, NE

A Concise Guide to Market Research

The Process, Data, and Methods Using IBM SPSS Statistics

Mooij, E.; Sarstedt, M.

2011, XX, 308 p. 130 illus., 49 illus. in color., Hardcover

ISBN: 978-3-642-12540-9